

Report TER

Julie Dornat

April 2026

1 Abstract

2 Introduction

Computational social choice [Brandt et al., 2016] studies collective decision-making problems involving multiple agents with potentially conflicting preferences. Its goal is to design mechanisms that aggregate individual preferences into a collective decision while satisfying desirable properties such as fairness and efficiency. Typical fairness criteria include maxmin and envy-freeness, which aim to capture the extent to which agents are satisfied with the outcome.

A central challenge in computational social choice lies in the elicitation and representation of agents' preferences over a set of possible alternatives or allocations of resources. Most existing results assume complete preference information, which limits their applicability in realistic scenarios. In many practical settings, it is unrealistic to assume that agents can fully specify their preferences, as the number of possible allocations grows combinatorially with the number of resources (*Handbook of Computational Social Choice, Chapter 10, p224*). This makes preference elicitation both cognitively demanding for agents and computationally challenging. As a result, a significant line of research focuses on developing compact and expressive preference models that mitigate these difficulties.

At the same time, collective decision-making must balance multiple, often conflicting objectives. As discussed in the *Handbook of Computational Social Choice, Chapter 12*, fairness criteria such as envy-freeness or maxmin may conflict with efficiency requirements such as Pareto efficiency or utilitarian optimality, which aim at maximizing the overall welfare and avoiding loss of resources. This results in seeking solution able to handle the tradeoff between fairness and efficiency.

This motivates the study of preference representations that capture only partial information, while still allowing one to reason about fairness and efficiency properties of the resulting outcomes. In this context, partial preferences provide a natural framework to model incomplete knowledge about agents' preferences. As emphasized by *Kenig and Kimelfeld* (2019), incomplete preferences constitute a distinct form of uncertainty.

Rather than assuming fully specified rankings, agents may only reveal limited or partial information. As discussed by *Jérôme Lang* (2020), this raises a fundamental question: how can one make reliable collective decisions when the available preference information is incomplete? and when can the eliciting process can be stopped to deduce reliable solution? To address this issue, the notions of possible and necessary outcomes have been introduced, allowing one to reason about decisions that hold under some or all completions of the missing information.

Computational social choice has applications in a wide range of domains, including voting, fair division, matching, multi-agent systems, and information retrieval. In this work, we focus in particular on fair division and matching problems.

Fair division concerns the allocation of resources among agents in a fair manner, and applies to both divisible and indivisible goods. A classical example in the indivisible setting is the house allocation problem, where each agent is assigned a single house. Such settings can also be viewed through the lens of matching theory, where agents and objects form two sets to be matched.

Matching theory provides a rich framework for modeling allocation mechanisms, including concepts such as perfect matching and popular matchings.

Research Question This report investigates the following question:

How can we design fair allocation or matching mechanisms when agents' preferences are only partially known?

More specifically, we explore how notions such as possible and necessary envy-freeness can be used to reason about fairness under incomplete preferences, and how these ideas can be connected to matching frameworks.

3 Preference Models

Partial preferences aim to represent situations in which agents do not fully specify their preferences over all possible alternatives. Instead of a complete description, only partial information is available, which induces uncertainty regarding the evaluation of outcomes with respect to properties such as fairness or efficiency.

From an informational perspective, partial preferences can be seen as a compact representation of agents' preferences. This is closely related to the notion of representation and compilation complexity discussed in the *Handbook of Computational Social Choice*, Chapter 10, section 10.7, where the goal is to encode preference information in a concise way while preserving the ability to reason about it.

During our literature review, we identified several families of partial preference models, which differ in the type and amount of information elicited from agents.

(1) Preference Representation Models

- **Ordinal preferences:** agents provide a ranking over alternatives or subsets of alternatives. Extensions such as the responsive set extension (*Fair Division under Ordinal Preferences: Computing Envy-Free Allocations of Indivisible Goods* Bouveret et al, Brams et al (2005)) allow one to derive preferences over bundles from rankings over individual items.
- **Cardinal preferences:** agents assign numerical utilities to alternatives, capturing the intensity of preferences.

(2) Partial Elicitation Models

- **Truncated preferences:** agents report only part of their ranking.
 - **Top- k preferences:** only the k most preferred alternatives are reported.
 - **Bottom- k preferences:** only the k least preferred alternatives are reported, leaving the rest unranked. In this case, the most preferred alternative may remain unknown.
 - **Set ranking queries:** agents provide a ranking into a subset of alternatives.
 - **Intervals of cardinal preferences:** agents provide ranges of utility values instead of a single utility.
- **Pairwise comparisons:** agents indicate which alternative is preferred between two options. These comparisons are central in majority-based methods and Condorcet-based reasoning (e.g., Condorcet winner, Copeland's rule).
- **Positional queries:** agent provides the alternative she ranks at a given position, leading to the construction of scoring vector. Using these scoring vectors in voting context, Kenig and Kimelfeld (2019) analyze the computational complexity of estimating the probability of an election outcome under uncertainty.
- **Approval and veto preferences:**
 - **Approval preferences:** agents classify alternatives into approved and non-approved sets.
 - **Veto preferences:** agents reject a subset of alternatives.

- **Acceptability-based preferences:** agents specify which alternatives are acceptable, often modeled using bipartite graphs in matching problems. For instance, in the marriage problem, some agents may find certain allocations unacceptable. Thus, not all objects are allocable to a given agent. Acceptability therefore amounts to defining a list of preferences by excluding those on which the agent issues a right of veto.
- **Query-based elicitation models:**
 - **Next-best queries:** agents iteratively reveal their most preferred alternative among those the remaining ones. We call it top k queries when we stop the retrieval at rank k.
 - **Split queries:** agents partition a subset of alternatives into ordered groups. More precisely, the Chapter 10 of the Handbook of Computational Social Choice explains that agent is asked to partition a subset of alternatives into two or more subsets and order them such that all alternatives in a higher-ranked subset are preferred to all alternatives in a lower-ranked subset. The k-approval is a special case where two blocks of size k and m-k are created.
 - **Choice queries:** agents select their most preferred alternative from a subset.

(3) Structured Preference Domains

- **Single-peaked preferences:** agents have a most preferred option, and their preferences decrease as they move away from this peak. It assumes that there exists an axis on which alternatives are positionned independently of the agents. In the paper "Single-peaked consistency and its complexity" Bruno Escoffier, Jérôme Lang, Meltem "Ozt"urk study whether the preference profiles are single peaked and how to determine the axis based on some agents' preferences.
- **Euclidean preferences:** preferences are derived from distances in a geometric space. This multidimensional space is defined by characteristics used to describe and compare alternatives. The paper "Euclidean preferences" of Anna Bogomolnaia and Jean-François Laslier defines Euclidean preferences, such as the geometric representation as a point in a d-dimensional space. The alternative is judged as good as it is close to the ideal point given by the best preference of an agent. The authors also consider convex preferences, which induce intermediate preferences between two options that agent prefers. They study propositions over the sufficient dimension for representing any Euclidean preference profile.
- **Attribute-based preferences:** agents have different preference profile depending on the attribute considered (politics, economics, etc).

(4) Probabilistic Models

- **Probabilistic preference models:** preferences are drawn from a probability distribution (e.g., impartial culture or Polya-Eggenberger urn models), allowing one to reason under uncertainty.

The Chapter 10 of the Handbook of Computational Social Choice, presents most of these preferences models.

These models illustrate different ways of reducing the amount of information required from agents while maintaining the ability to reason about collective decisions. Not all these models are equally relevant for the study of fairness under incomplete information, which motivates focusing on specific classes in the remainder of this work.

In practice, the choice of a preference model has a direct impact on both the expressiveness of the representation and the computational complexity of the associated problems. It therefore depends on the target application and the type of reasoning required. Partial preference models are widely used in applications such as voting and matching problems (e.g., house allocation, hospital-resident or marriage problems). Simpler models reduce the elicitation burden and facilitate computation, while richer models allow for more expressiveness at the cost of increased complexity.

A key advantage of partial preferences is that they reduce the amount of information agents need to communicate, which can improve the reliability of the elicited data while maintaining sufficient structure for decision-making.

4 Preliminaries

for all agents $i \in N$

5 Fair division with partial preferences

6 Matching and popularity

7 Possibly and Necessarily Popular Matching

8 Openings

References

Felix Brandt, Vincent Conitzer, Ulle Endriss, Jérôme Lang, and Ariel D Procaccia. *Handbook of computational social choice*. Cambridge University Press, 2016.